

Where Substitution Happens

Semantic substitution, deferred unquote and explicit substitution
in a paired-message rho calculus

L.G. Meredith
F1R3FLY.io

Working note, 1 October 2026

Abstract

The note *Unitary in the Equations, Measurement in the Reductions* [4] arrived, on the way to a quantum interpretation, at a small extension of the rho calculus that stands on its own: messages and names are pairs of a process and a datum drawn from a type D , dropping splits into a binary equation that preserves the datum and a unary drop resolved by semantic substitution, and replication is derived by reflection. This note extracts that core, independently of any quantum reading, and uses it to compare three treatments of information exchange at interaction. In the first, the one adopted, the received code is installed by semantic substitution at the communication itself. In the second, substitution is syntactic and a separate unary rule $\& @ \langle p, d \rangle \rightsquigarrow p$ unquotes later. In the third, the substitution is a process term left beside the continuation, in the style of explicit fusions. The two alternatives were proposed by Mike Stay. The argument against them has three parts: unary reductions that no interaction triggers do not fit the context-labelled bisimulation programme; both alternatives change the shape of bisimulation; and that change is the formal shadow of a practical cost, because once substitution is detached from communication the programmer must coordinate by hand which substitutions happen together and which happen in order. Worked examples show both burdens: on the programmer, and on the bisimulation proofs.

Contents

1	Introduction	2
2	The core calculus	2
3	The two alternatives	4
3.1	Deferred unquote	4
3.2	Explicit substitution	4
4	Where the unquote lives	5
5	The argument	5
6	The burden on the programmer	6
6.1	Knowing that received code is live	6
6.2	Substitutions that must happen in order	7
6.3	Substitutions that must happen together	7
6.4	Names computed from received code	7

7	The burden on bisimulation	8
7.1	Unary redexes and labels	8
7.2	The relations get larger	8
7.3	Laws that fail outright	9
7.4	Substitutions in the state	9
7.5	Cost	9
8	What the alternatives offer	9
9	Open questions	10

1 Introduction

Every process calculus must say what happens to the information carried by a message at the moment it is received. In the ρ calculus [5] the question has an extra edge, because what is received is a quoted process, and the receiver may do two quite different things with it: treat it as a name, or run it. The original paper answers with *semantic substitution*: substitution is syntactic everywhere except at a drop $*y$, where it installs the received process in place. The quantum design session recorded in [4] rediscovered this choice under pressure, and in doing so produced a version of the calculus in which the choice is unusually easy to see. Names and messages are pairs, so there is a datum to lose; the binary drop keeps it and the unary drop erases it; and nothing else in the calculus touches the question.

That makes the core calculus a good place to compare alternatives. In conversation, Mike Stay proposed two. The first keeps substitution uniformly syntactic and moves execution of received code into a separate unary reduction. The second makes substitution a process term, after Wischik and Gardner’s explicit fusions [2]. Both are natural, and both are standard moves elsewhere in the literature on substitution [1]. This note records why the paired-message calculus keeps semantic substitution, and makes the reasons concrete with examples.

Reading order. Section 2 states the core calculus. Section 3 states the two alternatives. Section 4 compares all three by asking where the unquote, and with it the erasure of the datum, takes place. Section 5 states the argument. Sections 6 and 7 give the examples, first the programmer’s burden and then the burden on bisimulation proofs. Section 8 records what the alternatives do offer, and Section 9 lists what remains open.

2 The core calculus

The construction is a type parameterisation of the ρ calculus. Fix a nonempty set D of data, and a distinguished $d_0 \in D$. Nothing below uses any structure on D ; in [4] it is the group $\mathbb{Z}/8$ of phase labels, but here it is arbitrary.

Definition 2.1 (Syntax of $\text{Sem}[D]$).

$$\begin{array}{ll}
M & ::= \langle P, d \rangle & \text{messages, } d \in D \\
x, u & ::= @M \mid y & \text{names} \\
P & ::= 0 \mid x!(M) \mid \text{for}(y \leftarrow x) P \mid P \mid P \mid *(x, u) \mid *y & \text{processes}
\end{array}$$

The receipt $\text{for}(y \leftarrow x) P$ binds y in P . The unary drop $*y$ applies only to variables. There is no restriction operator and no persistent receipt: fresh names arise by quotation, and replication is derived (Proposition 2.7).

Definition 2.2 (Structural congruence). $\equiv_S$ is the least congruence, closed under α -conversion, satisfying

(E1) Monoid

$$P \mid 0 \equiv_S P, \quad P \mid Q \equiv_S Q \mid P, \quad (P \mid Q) \mid R \equiv_S P \mid (Q \mid R).$$

(E2) Binary drop

$$*(\langle p, d \rangle, u) \equiv_S u!(\langle p, d \rangle).$$

Names are identified up to the congruence of what they quote: $\langle p, d \rangle = \langle q, e \rangle$ iff $p \equiv_S q$ and $d = e$.

Remark 2.3 (Why E2 is an equation). E2 could have been read as a notational abbreviation instead. It is folded into $\equiv_S$ for two reasons. As an equation it is a bijection between term shapes: every output $u!(M)$ is congruent to exactly one binary drop $*(\langle M, u \rangle)$, and nothing is lost in either direction. And because it sits in the quotient, the observation theory of Section 7, which is defined up to $\equiv_S$, never has to distinguish a forwarded message from its expansion. The binary drop forwards the whole pair; in particular it *preserves the datum*.

Definition 2.4 (Semantic substitution). Let $N = \langle p, d \rangle$. The substitution $\{N/y\}$ acts on names by $y\{N/y\} = N$, $z\{N/y\} = z$ for $z \neq y$, and $\langle Q, e \rangle\{N/y\} = \langle Q\{N/y\}, e \rangle$. On processes it commutes with every constructor, renaming bound variables to avoid capture, with one exception:

$$(*y)\{\langle p, d \rangle/y\} = p.$$

In particular $*(y, u)\{N/y\} = *(N, u) \equiv_S u!(\langle p, d \rangle)$ keeps the datum, while the unary drop discards it. Substitution acts under quotation, so a name built from received code changes at the moment the code arrives.

Definition 2.5 (Reduction).

$$\begin{aligned} (\text{COMM}) \quad & \text{for}(y \leftarrow x) P \mid x!(\langle q, d \rangle) \rightsquigarrow P\{\langle q, d \rangle/y\} \\ (\text{PAR}) \quad & P \rightsquigarrow P' \Rightarrow P \mid Q \rightsquigarrow P' \mid Q \\ (\text{STRUCT}) \quad & P \equiv_S P', P' \rightsquigarrow Q', Q' \equiv_S Q \Rightarrow P \rightsquigarrow Q \end{aligned}$$

COMM is the only rule that does any work, and every reduction is the meeting of a sender and a receiver.

Proposition 2.6 (Closure under reflection). *Every communication substitutes a name of the form $\langle M, u \rangle$, and E2 takes such a name back to a message. The payload sort and the quoted sort coincide.*

Proposition 2.7 (Replication is derivable). *Let $R_x := \text{for}(y \leftarrow x) (*y, x) \mid *y$ and $!P := x!(\langle R_x \mid P, d_0 \rangle) \mid R_x$. Then*

$$!P \rightsquigarrow x!(\langle R_x \mid P, d_0 \rangle) \mid R_x \mid P.$$

Proof. COMM substitutes $N = \langle R_x \mid P, d_0 \rangle$ for y . The binary drop becomes $*(N, x) \equiv_S x!(\langle R_x \mid P, d_0 \rangle)$, which re-emits the message with its datum intact, and the unary drop becomes $R_x \mid P$. $\square$

The proof shows the two drops doing different jobs within a single step: one forwards a copy for the next unfolding, the other installs a copy to run. Both happen at once, because both are effects of the same substitution.

Remark 2.8 (Installation and execution). Semantic substitution fixes the moment received code is *installed*: at the COMM. It does not force the code to run then. If $*y$ is guarded, say $\text{for}(y \leftarrow x) \text{for}(z \leftarrow c) *y$, then p is installed under the guard at the first COMM and runs when the guard is lifted by the second. Both moments are communications, and both are under the program's control.

3 The two alternatives

3.1 Deferred unquote

Definition 3.1 ($\text{Unq}[D]$). Replace the unary drop with an unquote operator, and substitute syntactically:

$$\begin{aligned} P &::= 0 \mid x!(M) \mid \text{for}(y \leftarrow x) P \mid P \mid P \mid *(x, u) \mid \& x \\ (\text{COMM}_U) & \quad \text{for}(y \leftarrow x) P \mid x!(\langle q, d \rangle) \rightsquigarrow P\{\@ \langle q, d \rangle / y\}_{\text{syn}} \\ (\text{UNQ}) & \quad \& \@ \langle p, d \rangle \rightsquigarrow p \end{aligned}$$

with E1, E2, PAR and STRUCT as before.

The operator must apply to arbitrary names, not only to variables, because syntactic substitution turns $\& y$ into $\& \@ \langle q, d \rangle$. A consequence is that $\& \@ \langle p, d \rangle$ can be written directly in source text, as a term that reduces with nobody else involved. In $\text{Sem}[D]$ the corresponding term is never formed. The binary drop is unaffected: substitution into $*(y, u)$ is already syntactic, and E2 applies.

3.2 Explicit substitution

The second proposal leaves the substitution itself in the term, more or less following explicit fusions [2, 6]. The rules below are a reconstruction, made to have something definite to argue about; the precise rules intended may differ, and the argument is written so that it does not depend on the details.

Definition 3.2 ($\text{ES}[D]$). Extend the syntax of $\text{Unq}[D]$ with a substitution term and a scope for the variable it governs:

$$\begin{aligned} P &::= \dots \mid [N/y] \mid (y) P \\ (\text{COMM}_E) & \quad \text{for}(y \leftarrow x) P \mid x!(M) \rightsquigarrow (y) (P \mid [@M/y]) \end{aligned}$$

with UNQ for unquoting, and the fusion-style equations

(X1) Propagation

$$(y) (P \mid [N/y]) \equiv_S (y) (P\{N/y\}_{\text{syn}} \mid [N/y]).$$

(X2) Collection

$$(y) [N/y] \equiv_S 0.$$

(X3) Extrusion

$$((y) P) \mid Q \equiv_S (y) (P \mid Q) \text{ when } y \notin \text{fv}(Q).$$

The scope is needed. In the receipt, y is bound by the input prefix; COMM_E removes the prefix, so something else must keep y from becoming a free, global variable. In the fusion calculus that job is done by restriction, which keeps fusions local. The ρ calculus has no restriction, deliberately, since fresh names arise by quotation, so $\text{ES}[D]$ must import one.

There are two ways to read the propagation of the substitution term, and they lead to different calculi.

- **Equational reading.** X1 holds as stated. Then by X1, X3 and X2, $(y) (P \mid [N/y]) \equiv_S P\{N/y\}_{\text{syn}}$, so $\text{ES}[D]$ coincides with $\text{Unq}[D]$ up to $\equiv_S$ and adds only the bookkeeping of the scope. Everything said about $\text{Unq}[D]$ below then applies to it unchanged.

- **Interactional reading.** $X1$ is dropped, and each unguarded occurrence of y (as a channel subject, or under $\&$) is resolved by its own reduction against the substitution term, which persists until collected: for instance $[@\langle p, d \rangle / y] \mid \& y \rightsquigarrow [@\langle p, d \rangle / y] \mid p$. Each resolution is now an interaction, which answers the first objection below, but a single COMM is followed by one resolution step per occurrence, and occurrences under quotation are never reached, because nothing reduces inside a name.

4 Where the unquote lives

The three calculi agree on everything except the treatment of the unary drop. The binary drop is the same equation in all of them. Without the unary drop, nothing runs received code, recursion is lost, and with linear receipts every term strongly normalises [4]. So the unary drop is exactly where recursion comes from, and the whole comparison concerns one operator.

The unary drop is also the one place the datum is erased: $@\langle p, d \rangle \mapsto p$ is many-to-one. It is therefore useful to ask, of each design, which step performs that erasure. A fourth placement, the equation $*@\langle p, d \rangle \equiv_S p$, completes the picture.

Placement	Triggered by interaction	Atomic with the COMM	Where d is erased
Equation, $*@\langle p, d \rangle \equiv_S p$	no step	no step	in $\equiv_S$, at no cost
Semantic substitution ($\text{Sem}[D]$)	yes	yes	at the COMM
Deferred unquote ($\text{Unq}[D]$)	no	no	at a UNQ step, per occurrence
Explicit substitution ($\text{ES}[D]$), equational	no	no	as in $\text{Unq}[D]$
Explicit substitution ($\text{ES}[D]$), interactional	yes	no	at each resolution

The equational placement is ruled out in [4] because equations are free while information erasure is not: an equation that forgets d makes $\equiv_S$ identify configurations that differ in data, and cost accounting has nowhere to charge for the loss. Semantic substitution charges the erasure to the COMM, which is already the step that interaction pays for. The deferred unquote charges it to a step that no interaction pays for. Explicit substitution either inherits that placement or spreads the erasure across as many steps as there are occurrences.

5 The argument

1. **Reductions should be interactions.** A unary rule like $\& @\langle p, d \rangle \rightsquigarrow p$, which fires without any interaction triggering it, does not fit the context-labelled bisimulation programme [3]. There, a transition is labelled by the minimal context that completes a redex, and every redex of $\text{Sem}[D]$ is a sender and receiver meeting. A redex that is complete on its own contributes transitions that no context takes part in, and that the environment can neither cause nor see except through their effects.
2. **Both proposals change bisimulation.** $\text{Unq}[D]$ adds internal steps between a communication and the running of what it delivered. $\text{ES}[D]$ puts substitution terms into the states being compared. Either way the equivalence, and the proofs about it, must be rethought.
3. **That change reflects a cost to the programmer.** In $\text{Sem}[D]$ the substitution happens atomically, at the COMM, and the programmer coordinates when and where substitutions

happen with the full power of the language, because every substitution is a communication. In the alternatives, substitutions are detached from the communications that produced them, and the programmer has to do extra work to make sure that substitutions that must coincide do, and substitutions that must be ordered are.

Points 2 and 3 are one point seen from two sides: the more complicated bisimulation is the formal record of the coordination the programmer now has to do. The next two sections give examples of each side.

6 The burden on the programmer

Throughout, a, b, c, x, x_1, x_2 and ack are fixed, distinct quoted names, and $e \in D$ is an arbitrary datum. Write $P \downarrow_a$ when $P \equiv_S a!(M) \mid Q$ for some M and Q , so that P offers an output on a .

6.1 Knowing that received code is live

Example 6.1 (Install, then acknowledge). A receiver installs code it is sent and then acknowledges:

$$S = x!(\langle p, d \rangle) \mid \text{for}(y \leftarrow x) (*y \mid ack!(\langle 0, e \rangle)).$$

In $\text{Sem}[D]$, $S \rightsquigarrow p \mid ack!(\langle 0, e \rangle)$. Whoever consumes the acknowledgement does so in a state where p is already installed: the acknowledgement means what it says. In $\text{Unq}[D]$,

$$S \rightsquigarrow \& @ \langle p, d \rangle \mid ack!(\langle 0, e \rangle),$$

and the acknowledgement can be consumed, and acted on, before UNQ fires.

The natural repair is for the receiver to acknowledge *after* the unquote. That cannot be done from the receiver's side. Suppose the receiver wraps the code, sends the wrapper to itself, and unquotes it:

$$\text{for}(y \leftarrow x) (w!(\langle \& y \mid ack!(\langle 0, e \rangle), e \rangle) \mid \text{for}(z \leftarrow w) \& z).$$

After both communications and the UNQ of the wrapper, the state is $\& @ \langle p, d \rangle \mid ack!(\langle 0, e \rangle)$, which is where it started. Every wrapper leaves the inner unquote pending.

Proposition 6.2 (No receiver-side signal of liveness). *In $\text{Unq}[D]$, no part of the receiver can be causally downstream of the step $\& @ \langle p, d \rangle \rightsquigarrow p$, except through actions of p itself.*

Proof sketch. UNQ has a single residual, p , and no partner. Anything else in the term is unchanged by the step and so cannot depend on it. The only events that follow it causally are those of p . $\square$

So in $\text{Unq}[D]$ the acknowledgement must be built into the shipped code: the *sender* must send $\langle p \mid ack!(\langle 0, e \rangle), d \rangle$. That moves a receiver's coordination problem across the trust boundary to the sender, and it changes the message, hence the name $@M$, hence any protocol keyed on that name. In $\text{Sem}[D]$ the receiver needs no help. Its continuation running is the evidence that the code is installed, because both are the same step.

6.2 Substitutions that must happen in order

Example 6.3 (Activation follows reception). A receiver takes two pieces of code in sequence and runs each:

$$T = x!(\langle p_1, d_1 \rangle) \mid x!(\langle p_2, d_2 \rangle) \mid \text{for}(y_1 \leftarrow x) (*y_1 \mid \text{for}(y_2 \leftarrow x) *y_2).$$

Taking p_1 as the first one received, in $\text{Sem}[D]$ every run passes through $p_1 \mid \text{for}(y_2 \leftarrow x) *y_2 \mid x!(\langle p_2, d_2 \rangle)$: the first code is installed before the second can be received. In $\text{Unq}[D]$ this run is available:

$$\begin{aligned} T &\rightsquigarrow \& @\langle p_1, d_1 \rangle \mid \text{for}(y_2 \leftarrow x) \& y_2 \mid x!(\langle p_2, d_2 \rangle) \\ &\rightsquigarrow \& @\langle p_1, d_1 \rangle \mid \& @\langle p_2, d_2 \rangle \\ &\rightsquigarrow \& @\langle p_1, d_1 \rangle \mid p_2 \\ &\rightsquigarrow p_1 \mid p_2. \end{aligned}$$

The second code is live while the first is not.

The order of reception no longer fixes the order of activation. When that order matters, as when p_2 is an update meant to supersede p_1 , the programmer must restore it, and by Proposition 6.2 the only way is the one from Example 6.1: the sender of p_1 builds in a signal, and the receiver guards the second receipt on it.

6.3 Substitutions that must happen together

Example 6.4 (Simultaneous availability). Two pieces of code arrive on different channels, from different senders, and should become available together. Let $p_1 = a!(\langle 0, e \rangle)$ and $p_2 = b!(\langle 0, e \rangle)$, and

$$U = x_1!(\langle p_1, d_1 \rangle) \mid x_2!(\langle p_2, d_2 \rangle) \mid \text{for}(y_1 \leftarrow x_1) \text{for}(y_2 \leftarrow x_2) (*y_1 \mid *y_2).$$

In $\text{Sem}[D]$, p_1 is installed under the guard on x_2 at the first COMM, and both are released by the second: $U \rightsquigarrow \rightsquigarrow p_1 \mid p_2$. Along every run, a state offers output on a only if it also offers output on b : $P \downarrow_a$ implies $P \downarrow_b$. In $\text{Unq}[D]$, after both COMMs and one UNQ the state is $p_1 \mid \& @\langle p_2, d_2 \rangle$, with $\downarrow_a$ but $\not\downarrow_b$.

As in the previous examples, the receiver cannot repair this. Packaging the two received codes into one wrapper, $w!(\langle \& y_1 \mid \& y_2, e \rangle)$, and unquoting it leaves $\& @\langle p_1, d_1 \rangle \mid \& @\langle p_2, d_2 \rangle$, two independent pending steps again. In $\text{Unq}[D]$, two codes received from independent senders cannot be made live simultaneously by the receiver alone. In $\text{Sem}[D]$ it takes no effort at all, because a single COMM installs every drop of the variables it has bound so far.

With k ungarded drops of received variables, a single COMM of $\text{Sem}[D]$ corresponds in $\text{Unq}[D]$ to a COMM followed by k independent unquotes, with 2^k reachable states of partial activation. Each of these is a state the programmer must either show to be harmless or rule out.

6.4 Names computed from received code

Because substitution in $\text{Sem}[D]$ acts under quotation, a name built from received code is resolved at the COMM along with everything else. This is how `rho` programs compare code: two processes rendezvous on a name exactly when the code they quote is congruent.

Example 6.5 (Naming the code that arrived). Let $n(y) = @(*y, d_0)$, a name computed from received code, and

$$V = x!(\langle p, d \rangle) \mid \text{for}(y \leftarrow x) n(y)!(\langle 0, e \rangle) \mid \text{for}(z \leftarrow @\langle p, d_0 \rangle) Q.$$

In $\mathbf{Sem}[D]$, COMM gives $n(y)\{\textcircled{\langle p, d \rangle}/y\} = \textcircled{\langle p, d_0 \rangle}$, so the output meets the third component and Q runs. In $\mathbf{Unq}[D]$ the output lands on $\textcircled{\langle \& \textcircled{\langle p, d \rangle}, d_0 \rangle}$, which is a different name from $\textcircled{\langle p, d_0 \rangle}$, because UNQ is a reduction, not an equation, and nothing reduces inside a quote. The rendezvous never happens.

The workaround in $\mathbf{Unq}[D]$ is to listen on $\textcircled{\langle \& \textcircled{\langle p, d \rangle}, d_0 \rangle}$. But that name depends on the datum d of the message in which the code happened to travel. The same code obtained by three routes, received with datum d , received with datum d' , or written literally, gets three different names. Code identity becomes route-dependent, and the programmer must track the route of every piece of code whose name is used.

$\mathbf{ES}[D]$ under the interactional reading adds a variant of this burden even for channel positions. A standard ρ idiom is to reply on the name of the request itself:

$$x!(M) \mid \text{for}(r \leftarrow \textcircled{M}) S \mid \text{for}(y \leftarrow x) y!(\langle q, e \rangle).$$

In $\mathbf{Sem}[D]$ the server's reply is on $\textcircled{M}$ as soon as the request is received. Under the interactional reading the reply sits on the variable y until a separate resolution step against $[\textcircled{M}/y]$, and if the server replies twice, the two resolutions are independent and the simultaneity problem of Example 6.4 returns.

7 The burden on bisimulation

7.1 Unary redexes and labels

In the approach of Leifer and Milner [3], an agent a has a transition $a \xrightarrow{C} a'$ when $C[a] \rightsquigarrow a'$ and C is a minimal context that makes this happen; their congruence theorem is for the strong equivalences built on these labels. In $\mathbf{Sem}[D]$ every redex consists of a receiver and a sender, so a label is either the missing half of a communication or the identity, when both halves are already present in the agent. In $\mathbf{Unq}[D]$ the term $\& \textcircled{\langle p, d \rangle}$ is a complete redex on its own. It contributes identity-labelled transitions that correspond to no communication at all.

Strong equivalences count these steps. $\& \textcircled{\langle p, d \rangle}$ has exactly one transition and its target is p , so it is not strongly bisimilar to p except in degenerate cases. Every law of $\mathbf{Sem}[D]$ that involves running received code therefore survives in $\mathbf{Unq}[D]$ at best as a weak law, and extending the context-labelled congruence results from strong to weak equivalences requires further conditions that must then be checked.

7.2 The relations get larger

Consider a receiver with k unguarded drops of the received variable,

$$W = x!(\langle p, d \rangle) \mid \text{for}(y \leftarrow x) (*y \mid \cdots \mid *y).$$

In $\mathbf{Sem}[D]$, W has one reduction, to $p^{|k|}$, and any bisimulation containing W needs one more pair to account for it. In $\mathbf{Unq}[D]$, W reduces to $(\& \textcircled{\langle p, d \rangle})^{|k|}$, and then through every subset of the k unquotes. A weak bisimulation relating W to its $\mathbf{Sem}[D]$ counterpart must contain a pair for each of the 2^k partial states, each related to $p^{|k|}$ or to an intermediate of p 's own behaviour. These are the same 2^k states that burden the programmer in Example 6.4, now appearing as proof obligations. Up-to techniques can compress them, but they have to be developed and justified, whereas in $\mathbf{Sem}[D]$ they are not needed.

7.3 Laws that fail outright

The extra steps of $\text{Unq}[D]$ can be absorbed by moving to weak equivalence. The name problem of Example 6.5 cannot.

Proposition 7.1. *Let*

$$S_1 = x!(\langle p, d \rangle) \mid \text{for}(y \leftarrow x) @\langle *y, d_0 \rangle!(\langle 0, e \rangle),$$

$$S_2 = x!(\langle p, d \rangle) \mid \text{for}(y \leftarrow x) @\langle p, d_0 \rangle!(\langle 0, e \rangle),$$

with $y \notin \text{fv}(p)$, and $*y$ read as $\&y$ in $\text{Unq}[D]$. Then S_1 and S_2 are strongly barbed bisimilar in $\text{Sem}[D]$, and not even weakly barbed bisimilar in $\text{Unq}[D]$.

Proof sketch. In $\text{Sem}[D]$ both have the barb $\downarrow_x$ and a single reduction, to the same term $@\langle p, d_0 \rangle!(\langle 0, e \rangle)$. In $\text{Unq}[D]$ the reducts are $@(\& @\langle p, d \rangle, d_0)!(\langle 0, e \rangle)$ and $@\langle p, d_0 \rangle!(\langle 0, e \rangle)$. Both are stuck, and their barbs are on distinct names, since $\& @\langle p, d \rangle \not\equiv_S p$. No sequence of internal steps reconciles them, because no step acts inside a quote. $\square$

The bisimilarity in $\text{Sem}[D]$ is not a congruence: a context can send different code on x . The point is that in $\text{Unq}[D]$ the two terms are distinguished even without a context. The law “receiving code and naming it is the same as naming the code” holds in $\text{Sem}[D]$ and fails in $\text{Unq}[D]$, however weak the equivalence. Bisimulation proofs in $\text{Unq}[D]$ must keep track of the route by which each named piece of code arrived, which is Example 6.5 restated as a proof obligation.

7.4 Substitutions in the state

Under the equational reading, $\text{ES}[D]$ reduces to $\text{Unq}[D]$ up to $\equiv_S$, so bisimulation up to $\equiv_S$ absorbs the substitution terms and the previous subsections apply unchanged. Under the interactional reading the substitution terms are genuine components of the state. A bisimulation must relate $(y) (P \mid [N/y])$, and each of its partial resolutions, to $P\{N/y\}$, and the number of internal steps a single COMM gives rise to depends on the number of occurrences of y , which is a syntactic accident. Names with a free variable, such as $@(\&y, d_0)$, mean different channels under different surrounding substitutions, so barbs must be read relative to their scope. This is why the fusion calculus works with hyperequivalence, bisimulation closed under substitutions [6]. That is a heavier notion than anything $\text{Sem}[D]$ needs, since in $\text{Sem}[D]$ no state ever contains an unresolved received name.

7.5 Cost

In cost-accounted ρ , cost is charged to reductions. In $\text{Sem}[D]$ the cost of a run is a function of its communications. In $\text{Unq}[D]$ it also depends on how many drops the receiver happened to write, and under the interactional reading of $\text{ES}[D]$ on how many occurrences of the variable there are in any position. A cost-aware equivalence therefore distinguishes programs that differ only in how often they mention what they received. This is the cost-accounting version of the remark in Section 4: the alternatives charge for erasure at steps that interaction does not pay for.

8 What the alternatives offer

The alternatives have real merits, and the argument does not depend on denying them.

- **Provenance.** In $\text{Unq}[D]$ a name built from received code records the datum of the message the code arrived in. Sometimes that is wanted. In $\text{Sem}[D]$ it is available on request rather than imposed: the binary drop keeps the datum, so $@(* (y, u), d_0)$, which is congruent to $@(u!((p, d)), d_0)$ after the COMM, is a name that records d . The programmer chooses whether a name records where the code came from.
- **Sharing.** Explicit substitution earns its place in implementations, where it avoids copying large terms and lets a substitution be shared [1]. An implementation of $\text{Sem}[D]$ may well use explicit substitutions internally. The argument here is about the calculus, which fixes what is observable: whatever the machinery, the installation of received code must look atomic with the communication that delivered it.
- **Uniform substitution.** $\text{Unq}[D]$ has a simpler substitution operation, which is pleasant metatheoretically. The price is paid elsewhere: in an extra reduction rule, in weak equivalences, and in the failure of Proposition 7.1.

9 Open questions

Open question 9.1 (Operational correspondence). Translate $\text{Sem}[D]$ into $\text{Unq}[D]$ by reading $*y$ as $\&y$. Proposition 7.1 shows the translation is not adequate in general. Is it adequate up to weak bisimilarity on the fragment in which no drop occurs under quotation? The examples in Section 6 suggest it is, but this has not been proved.

Open question 9.2 (Relative pushouts with E2). With E2 in $\equiv_S$, every output is congruent to a binary drop. Labels can then be taken in a normal form with binary drops on quoted names expanded. It remains to check that the reactive system for $\text{Sem}[D]$ has the relative pushouts that the congruence theorem of [3] requires, and to state how semantic substitution interacts with the contexts involved.

Open question 9.3 (The explicit-substitution rules). The rules for $\text{ES}[D]$ here are a reconstruction. With the intended rules in hand, it should be checked which reading they fall under, and whether a middle reading exists that both keeps every resolution an interaction and restores atomicity.

Open question 9.4 (Structure on D). Nothing here uses structure on D . The quantum note [4] uses a group; other choices (costs, provenance tags, types) suggest themselves. The question is which properties of the core are stable under which choices, in particular whether erasure at the unary drop should be parametric in a map $D \rightarrow 1$ that the choice of D can refine.

Acknowledgement. The two alternatives, and the conversation that sharpened the argument, are due to Mike Stay. The note was drafted in dialogue with Claude (Anthropic).

References

- [1] M. Abadi, L. Cardelli, P.-L. Curien and J.-J. Lévy. Explicit substitutions. *J. Funct. Program.* 1(4):375–416, 1991.
- [2] P. Gardner and L. Wischik. Explicit fusions. *MFCS 2000*, LNCS 1893, 2000.
- [3] J. J. Leifer and R. Milner. Deriving bisimulation congruences for reactive systems. *CONCUR 2000*, LNCS 1877, 2000.
- [4] L. G. Meredith. Unitary in the equations, measurement in the reductions. Working note, F1R3FLY-io/publications, [fuzzyware/](https://fuzzyware.com/), 2026.

- [5] L. G. Meredith and M. Radestock. A reflective higher-order calculus. *ENTCS* 141(5):49–67, 2005.
- [6] J. Parrow and B. Victor. The fusion calculus: expressiveness and symmetry in mobile processes. *LICS 1998*.